

		ALGORITHM	DESCRIPTION	APPLICATIONS	ADVANTAGES	DISADVANTAGES
Supervised Learning	Linear Models	Linear Regression	A simple algorithm that models a linear relationship between inputs and a continuous numerical output variable	USE CASES <ol style="list-style-type: none"> Stock price prediction Predicting housing prices Predicting customer lifetime value 	<ol style="list-style-type: none"> Explainable method Interpretable results by its output coefficients Faster to train than other machine learning models 	<ol style="list-style-type: none"> Assumes linearity between inputs and output Sensitive to outliers Can underfit with small, high-dimensional data
		Logistic Regression	A simple algorithm that models a linear relationship between inputs and a categorical output (1 or 0)	USE CASES <ol style="list-style-type: none"> Credit risk score prediction Customer churn prediction 	<ol style="list-style-type: none"> Interpretable and explainable Less prone to overfitting when using regularization Applicable for multi-class predictions 	<ol style="list-style-type: none"> Assumes linearity between inputs and outputs Can overfit with small, high-dimensional data
		Ridge Regression	Part of the regression family — it penalizes features that have low predictive outcomes by shrinking their coefficients closer to zero. Can be used for classification or regression	USE CASES <ol style="list-style-type: none"> Predictive maintenance for automobiles Sales revenue prediction 	<ol style="list-style-type: none"> Less prone to overfitting Best suited where data suffer from multicollinearity Explainable & interpretable 	<ol style="list-style-type: none"> All the predictors are kept in the final model Doesn't perform feature selection
		Lasso Regression	Part of the regression family — it penalizes features that have low predictive outcomes by shrinking their coefficients to zero. Can be used for classification or regression	USE CASES <ol style="list-style-type: none"> Predicting housing prices Predicting clinical outcomes based on health data 	<ol style="list-style-type: none"> Less prone to overfitting Can handle high-dimensional data No need for feature selection 	<ol style="list-style-type: none"> Can lead to poor interpretability as it can keep highly correlated variables
	Tree-Based Models	Decision Tree	Decision Tree models make decision rules on the features to produce predictions. It can be used for classification or regression	USE CASES <ol style="list-style-type: none"> Customer churn prediction Credit score modeling Disease prediction 	<ol style="list-style-type: none"> Explainable and interpretable Can handle missing values 	<ol style="list-style-type: none"> Prone to overfitting Sensitive to outliers
		Random Forests	An ensemble learning method that combines the output of multiple decision trees	USE CASES <ol style="list-style-type: none"> Credit score modeling Predicting housing prices 	<ol style="list-style-type: none"> Reduces overfitting Higher accuracy compared to other models 	<ol style="list-style-type: none"> Training complexity can be high Not very interpretable
		Gradient Boosting Regression	Gradient Boosting Regression employs boosting to make predictive models from an ensemble of weak predictive learners	USE CASES <ol style="list-style-type: none"> Predicting car emissions Predicting ride hailing fare amount 	<ol style="list-style-type: none"> Better accuracy compared to other regression models It can handle multicollinearity It can handle non-linear relationships 	<ol style="list-style-type: none"> Sensitive to outliers and can therefore cause overfitting Computationally expensive and has high complexity
		XGBoost	Gradient Boosting algorithm that is efficient & flexible. Can be used for both classification and regression tasks	USE CASES <ol style="list-style-type: none"> Churn prediction Claims processing in insurance 	<ol style="list-style-type: none"> Provides accurate results Captures non linear relationships 	<ol style="list-style-type: none"> Hyperparameter tuning can be complex Does not perform well on sparse datasets
		LightGBM Regressor	A gradient boosting framework that is designed to be more efficient than other implementations	USE CASES <ol style="list-style-type: none"> Predicting flight time for airlines Predicting cholesterol levels based on health data 	<ol style="list-style-type: none"> Can handle large amounts of data Computational efficient & fast training speed Low memory usage 	<ol style="list-style-type: none"> Can overfit due to leaf-wise splitting and high sensitivity Hyperparameter tuning can be complex
	Clustering	K-Means	K-Means is the most widely used clustering approach—it determines K clusters based on euclidean distances	USE CASES <ol style="list-style-type: none"> Customer segmentation Recommendation systems 	<ol style="list-style-type: none"> Scales to large datasets Simple to implement and interpret Results in tight clusters 	<ol style="list-style-type: none"> Requires the expected number of clusters from the beginning Has troubles with varying cluster sizes and densities
		Hierarchical Clustering	A "bottom-up" approach where each data point is treated as its own cluster—and then the closest two clusters are merged together iteratively	USE CASES <ol style="list-style-type: none"> Fraud detection Document clustering based on similarity 	<ol style="list-style-type: none"> There is no need to specify the number of clusters The resulting dendrogram is informative 	<ol style="list-style-type: none"> Doesn't always result in the best clustering Not suitable for large datasets due to high complexity
		Gaussian Mixture Models	A probabilistic model for modeling normally distributed clusters within a dataset	USE CASES <ol style="list-style-type: none"> Customer segmentation Recommendation systems 	<ol style="list-style-type: none"> Computes a probability for an observation belonging to a cluster Can identify overlapping clusters More accurate results compared to K-means 	<ol style="list-style-type: none"> Requires complex tuning Requires setting the number of expected mixture components or clusters
Unsupervised Learning	Association	Apriori algorithm	Rule based approach that identifies the most frequent itemset in a given dataset where prior knowledge of frequent itemset properties is used	USE CASES <ol style="list-style-type: none"> Product placements Recommendation engines Promotion optimization 	<ol style="list-style-type: none"> Results are intuitive and Interpretable Exhaustive approach as it finds all rules based on the confidence and support 	<ol style="list-style-type: none"> Generates many uninteresting itemsets Computationally and memory intensive. Results in many overlapping item sets